


---

# 0. Erlang

Sequential Programming

Facultat d'Informàtica de Barcelona (FIB)  
Universitat Politècnica de Catalunya (UPC)



# Why Erlang?

---

- We need a language to implement some of the algorithms that we will learn
- Erlang is a general-purpose language and runtime environment well suited for scalable concurrent and distributed programming
  - Developed at Ericsson in late eighties
  - Functional Programming (e.g. Prolog, Lisp)
  - Available as free and open-source software
  - Built-in support for concurrency, distribution, and fault tolerance

# Why Erlang?

---

- Companies use Erlang in their production systems, e.g. Amazon, Yahoo!, Facebook, WhatsApp, T-Mobile, Motorola, Ericsson, ...
  - [https://en.wikipedia.org/w/index.php?title=Erlang\\_\(programming\\_language\)&oldid=754567048#Companies\\_using\\_Erlang](https://en.wikipedia.org/w/index.php?title=Erlang_(programming_language)&oldid=754567048#Companies_using_Erlang)
- Popular applications use Erlang, e.g. Chef, Ejabberd, CouchDB, GitHub, RabbitMQ, ...
  - [https://en.wikipedia.org/w/index.php?title=Erlang\\_\(programming\\_language\)&oldid=754567048#Software\\_projects\\_written\\_in\\_Erlang](https://en.wikipedia.org/w/index.php?title=Erlang_(programming_language)&oldid=754567048#Software_projects_written_in_Erlang)

# Relevant primitive data types

---

- atom: literal, constant with name (foo, my\_F, ...)
  - To be enclosed in single quotes (') if it does not begin with a lower-case letter or if it contains other characters than alphanumeric characters, '\_', or '@'
    - 'Monday', 'phone number', ...
  - Atoms true and false used to denote Boolean values
    - operators: not, and, or, xor, andalso, orelse
- number
  - integers: 10, -24, 16#A42B (i.e., base#value), \$a (\$char, i.e., ASCII code of the character char), ...
  - floats: 17.3, -56.62, 2.3e3, ...
  - ops: +, -, \*, /, div, rem, bnot, band, bor, bxor, bsl, bsr

# Relevant primitive data types

---

- reference: globally unique term
  - > **make\_ref()**.
  - #Ref<0.2853579700.305135618.16717>
- fun: to create an anonymous function
  - Can be passed as argument to higher-order functions
  - > F = **fun (X) -> X+1 end**.
  - #Fun<erl\_eval.6.39074546>
  - > F(2).
  - 3
- pid: process identifier (PID)
  - **self()** returns the PID of the calling process: <0.76.0>

# Relevant compound data types

---

- tuple: used to store a **fixed** number of terms
  - {123, def, abc}
  - {person, 'Joe', 'Armstrong'}
  - There exist a number of BIFs to manipulate tuples
- list: used to store a **variable** number of terms
  - [] (empty list)
  - [foo, 12, bar, zot]
  - See lists module for list processing functions
- "... " is shorthand for the list of ASCII codes of the enclosed chars within the quotes
  - "abcdefghi" is [97,98,99,100,101,102,103,104,105]

# Term comparison

---

- Operators:  $=$ ,  $<$ ,  $>$ ,  $=$ ,  $>$ ,  $==$  (equal to),  $===$  (exactly equal to: same value and type),  $\neq$  (not equal to),  $\neq$  (exactly not equal to)
- The following order is used to compare terms of different data types
  - *number < atom < reference < fun < port < pid < tuple < map < nil < list < bit string*
  - Lists compare element by element
  - Tuples compare by size and then element by element

$1 == 1.0$	$1 === 1.0$	$1 > ab$	$[1,2] < [2,1]$	$\{1,1\} < \{2\}$
true	false	false	true	false

# Variables

---

- Used to store values of terms
- The scope for a variable is its function clause
- Variables can only be bound **once**
  - The value of a variable can never be changed
- They start with an upper-case letter (or `'_'`)
  - `Abc`, `A_var`, `_Foo`
  - Variables starting with `'_'`, do not generate warnings when unused
  - The **anonymous variable** `'_'` can be used when a variable is required but its value can be ignored



# Pattern matching

---

- Assign values using pattern matching
  - $A = 10 \rightarrow$  Succeeds: binds A to 10
  - $B = \{z, \text{foo}, 4\} \rightarrow$  Succeeds: binds B to  $\{z, \text{foo}, 4\}$
  - $\{B, C, D\} = \{10, \text{foo}, \text{bar}\}$ 
    - Succeeds: binds B to 10, C to foo, D to bar
  - $\{A, A, B\} = \{\text{abc}, \text{abc}, \text{foo}\}$ 
    - Succeeds: binds A to abc, B to foo
  - $\{A, A, B\} = \{\text{abc}, \text{def}, 123\} \rightarrow$  Fails
  - $[A, B, C] = [1, 2, 3]$ 
    - Succeeds: binds A to 1, B to 2, C to 3
  - $[A, B, C, D] = [1, 2, 3] \rightarrow$  Fails

# Pattern matching

---

- Cons cell: **[ H | T ]**
  - Used for pattern matching on lists
  - The pattern "[H|T] = L" extracts the head into 'H' and tail into 'T' of the list 'L'
  - [H|T] = [1,2,3,4]
    - Succeeds: binds H = 1, T = [2,3,4]
  - [A,B|C] = [1,2,3,4,5,6,7]
    - Succeeds: binds A = 1, B = 2, C = [3,4,5,6,7]
  - [H|T] = [abc]
    - Succeeds: binds H = abc, T = []
  - [H|T] = [] → **Fails**

# Pattern matching

---

- Any element of a tuple or list can be of any type, even another tuple or list
  - $\{A, \_, [B | \_], \{B\}\} = \{abc, 23, [22, x], \{22\}\}$ 
    - Succeeds: binds  $A = abc$ ,  $B = 22$
- Pattern matching to get any element of a tuple
  - $Point = \{1, 4, 5\}$
  - $\{\_, Y, \_ \} = Point \rightarrow$  Succeeds: binds  $Y = 4$
  - $\{\_, Y\} = Point \rightarrow$  Fails
- Note the use of `'_'`, the anonymous variable, as a wildcard for pattern matching in both examples

# Function calls

---

- a) `Module:Function(Arg1, ..., ArgN)`
  - For external functions
    - `math2:double(10).`
    - `lists:keysearch(Name, 1, List).`
- b) `Function(Arg1, ..., ArgN)`
  - For local functions or auto-imported BIFs
    - `times(5, 2).`
    - `spawn(mod, init, []).`
  - `Arg1 ... ArgN` can be any Erlang terms
  - The module/function names must be atoms
    - Or expressions that evaluate to an atom

# Built In Functions (BIFs)

---

- `time()`
- `max(1, 2)`
- `length([1,2,3,4,5])`
- `is_tuple({a,b,c})`
- `size({a,b,c})`
- `element(2, {a,b,c})`
- `self()`
- `register(foo, Pid)`
- `link(Pid)`
- `make_ref()`
- The most commonly used BIFs belonging to the [erlang module](#) are auto-imported
- They do not need to be prefixed with the module name

# Function definition

---

```
func(Pattern11, Pattern21, ...) [when Guard1] -> ... ;  
func(Pattern21, Pattern22, ...) [when Guard2] -> ... ;  
...  
func(PatternN1, PatternN2, ...) [when GuardN] -> ... .
```

- Clauses are scanned sequentially until one matches its patterns with the given arguments and the guard, if any, is true
- Then, all variables in the heading become bound, the clause body is executed and the value of the last expression is returned
- Variables are local to each clause

# Function & module definition

---

- module(demo).
- export([double/1]).
- define(two, 2).
- double(X) -> times(X, ?two).
- times(X, N) -> X \* N.
- Functions are defined within Modules
  - Module name is to be same as the file name minus the .erl extension
- -export functions so that they can be called from outside the module: only double/1 is visible
  - double/1: the function 'double' with one argument
- -define macros and '?' to use them: ?two
  - Macros are expanded during compilation
  - There are some predefined macros: e.g., ?MODULE

# Function examples

---

```
-module(mathStuff).  
-export([factorial/1, area/1]).
```

```
factorial(0) -> 1;
```

```
factorial(N) -> N * factorial(N-1).
```

```
area({square, Side}) -> Side * Side;
```

```
area({circle, Radius}) -> 3.14 * Radius * Radius;
```

```
area({triangle, A, B, C}) -> S = (A + B + C)/2,  
    math:sqrt(S*(S-A)*(S-B)*(S-C));
```

```
area(Other) -> {invalid_object, Other}.
```



# Conditionals: If

---

if

Guard1 -> Body1;

... ;

GuardN -> BodyN

end.

fac(N) ->

if

N == 0 -> 1;

N > 0 -> N\*fac(N-1)

end.

- The branches are scanned sequentially until a guard sequence that evaluates to true is found
- If there is none, a run-time error occurs
- The guard expression true can be used in the last branch

# Conditionals: Case

---

```
case Expression of
  Pattern1 [when Guard1] -> ... ;
  ... ;
  PatternN [when GuardN] -> ...
end.
```

```
sum(L) ->
  case L of
    [] -> 0;
    [H|T] -> H + sum(T)
  end.
```

- Patterns are sequentially matched against the result of the expression
- If there is no matching pattern, a run-time error occurs

# Loops

---

- Functional programming languages usually do not offer looping constructs
- We must use recursion, notably, tail recursion
  - Used when the last expression of a function is a call to the same function
  - The stack frame of the current function is simply replaced with the one of the called function, allowing to implement loops efficiently in Erlang

```
loop(N) ->  
    io:format("~w~n", [N]),  
    loop(N+1).
```

# Writing output to terminal

---

- `io:format(Format, [Data])` BIF is used to write formatted data to the standard output
  - `Format`: string with formatting control sequences
    - `~B`: integer, `~f`: float (ddd.ddd), `~e`: float (d.ddde+-ddd)
    - `~g`: float (if  $\geq 0.1$  and  $< 10000.0$  as `~f`, else as `~e`)
    - `~s`: string, `~c`: ASCII code
    - `~w/~p`: any Erlang term (standard / pretty-printing)
    - `~n`: (platform-specific) new line sequence

```
io:format("s:~s w:~w p:~p B:~B c:~c g:~g~n",  
["Hi", "Hi", "Hi", 72, 72, math:pi()]).
```

```
s:Hi w:[72,105] p:"Hi" B:72 c:H g:3.14159
```

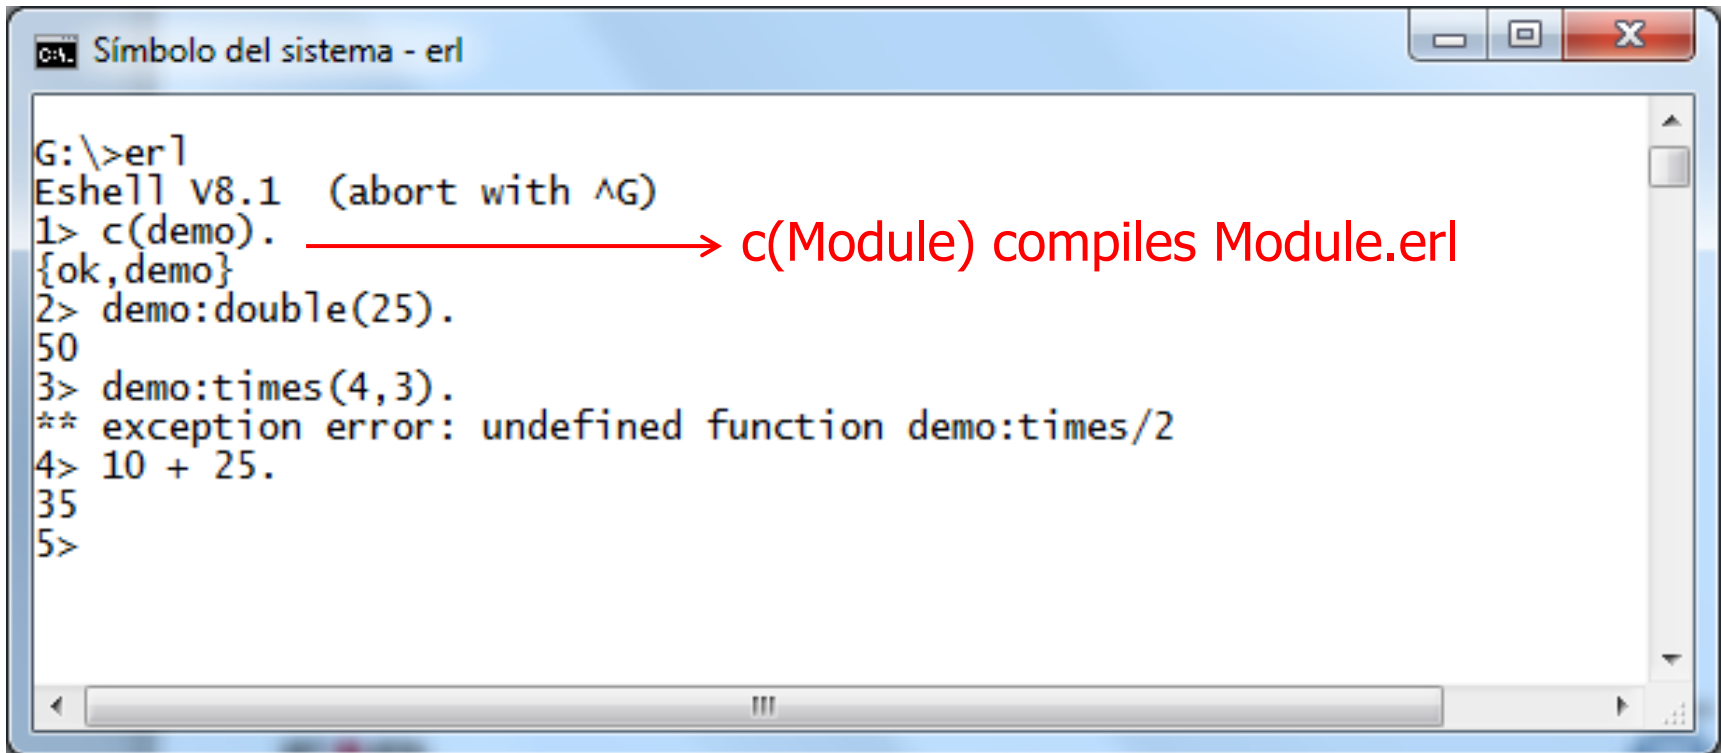
# Runtime system

---

- The Erlang runtime system gives you an interactive shell
- In the shell you can compile and load modules and call functions
- Run it by itself, inside vim, emacs, or in a IDE such as Eclipse

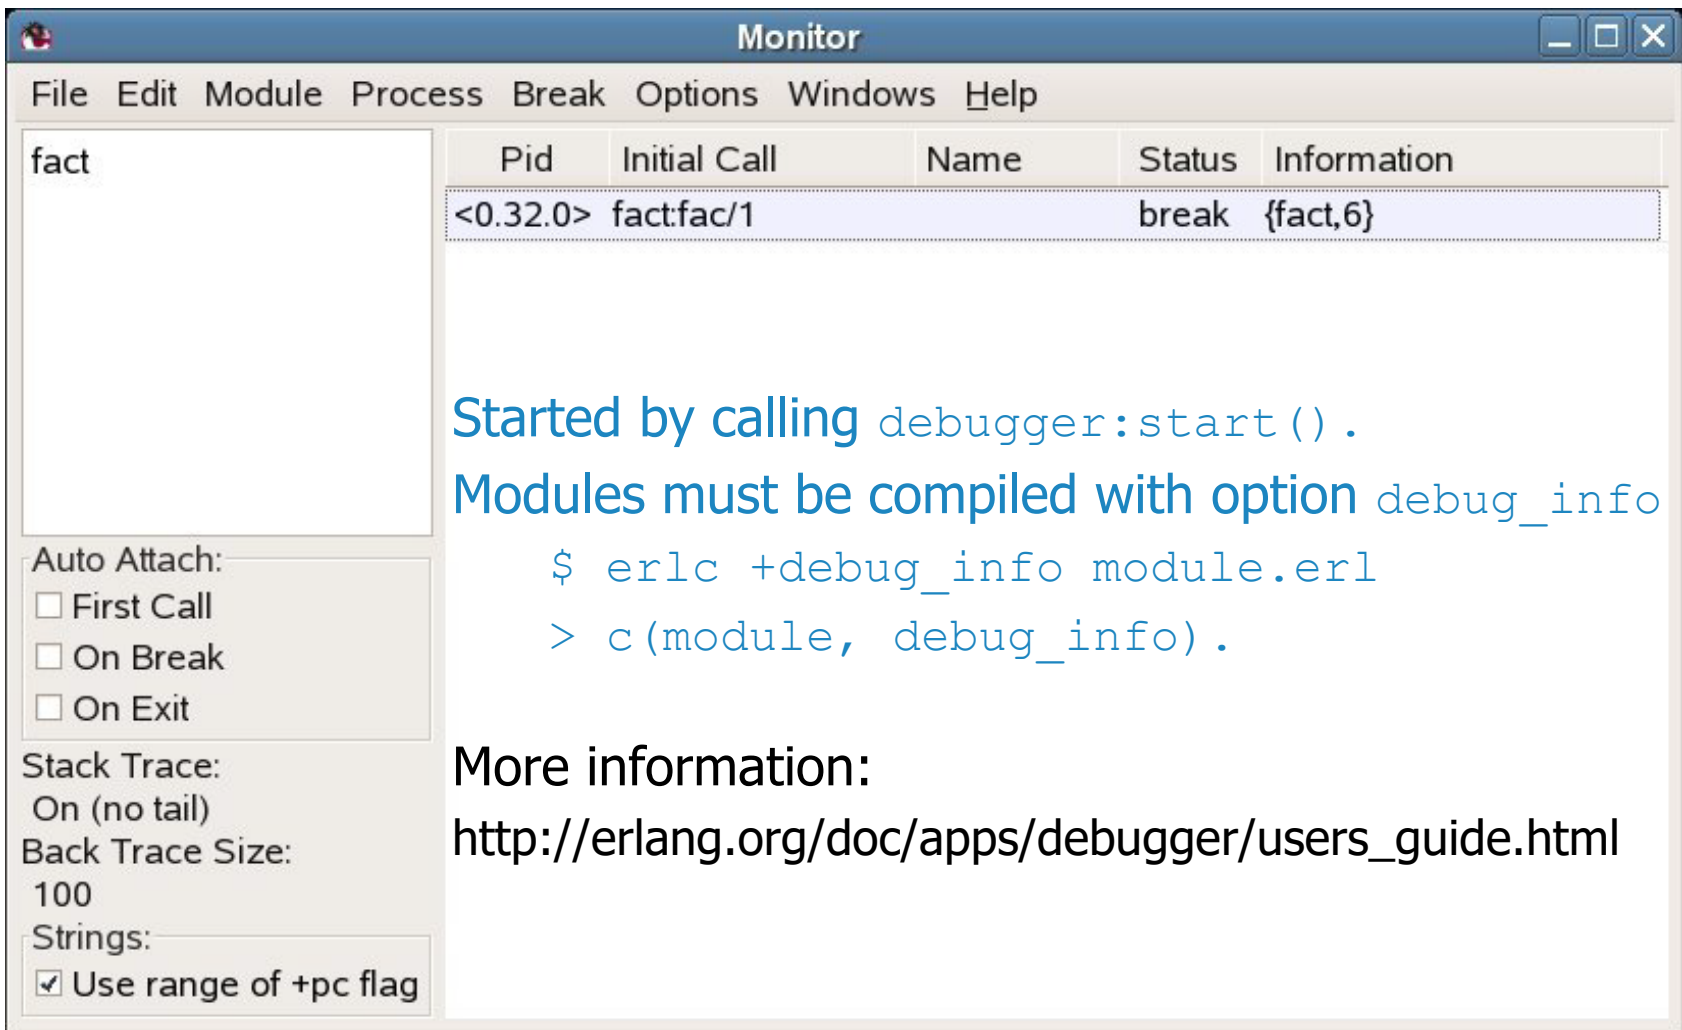
# Runtime system

---



```
Símbolo del sistema - erl
G:\>erl
Eshell V8.1 (abort with ^G)
1> c(demo). → c(Module) compiles Module.erl
{ok,demo}
2> demo:double(25).
50
3> demo:times(4,3).
** exception error: undefined function demo:times/2
4> 10 + 25.
35
5>
```

# Debugger



The screenshot shows the Erlang Monitor window with the following components:

- Menu Bar:** File, Edit, Module, Process, Break, Options, Windows, Help.
- Left Panel:**
  - fact:** A text area containing the word "fact".
  - Auto Attach:** Three checkboxes: ☐ First Call, ☐ On Break, ☐ On Exit.
  - Stack Trace:** On (no tail).
  - Back Trace Size:** 100.
  - Strings:** ☒ Use range of +pc flag.
- Table:**

Pid	Initial Call	Name	Status	Information
<0.32.0>	fact:fac/1		break	{fact,6}
- Main Area:**

Started by calling `debugger:start()`.

Modules must be compiled with option `debug_info`

```
$ erlc +debug_info module.erl
> c(module, debug_info).
```


More information:  
[http://erlang.org/doc/apps/debugger/users\\_guide.html](http://erlang.org/doc/apps/debugger/users_guide.html)

---

# 0. Erlang

Concurrent Programming

Facultat d'Informàtica de Barcelona (FIB)  
Universitat Politècnica de Catalunya (UPC)





# Process creation

---

- The Erlang BIF `spawn` is used to create and start the execution of a new process
  - a) `Pid = spawn(Module, Function, [Arg1, ..., ArgN])`
    - New process runs `Module:Function(Arg1,...,ArgN)`
    - Function must be exported from the module
  - b) `Pid = spawn(fun (Arg1, ..., ArgN) -> Body end)`
    - New process runs `Body`
- `spawn` returns the process identifier of the new process

# Message passing: send

---

- To send a message you need the process identifier of the receiver

Pid ! Message

- Message can be any valid Erlang terms
- The value of Message is also the return value of the expression
- Sending a message is **asynchronous**
  - No acknowledgement
- Sending a message to a pid never fails
  - Even if it refers to a non-existing process

# Message passing: receive

---

- All messages sent to a process are stored in its own queue in the order they are received
    - Messages from the same sender are ordered in the same order as they were sent (a.k.a. FIFO)
  - Selective receive: process can specify which messages it is willing to handle
  - Implicit deferral: messages remain in the queue until explicitly handled
- ↑ Can decide in what order to handle messages
- ↓ Queue may fill with forgotten messages

# Message passing: receive

---

- receive suspends a process waiting for a message that matches one of the patterns
  - receive matches the first message in the queue sequentially against the patterns
  - If no match is found, the procedure is repeated for the second message, and so on

receive

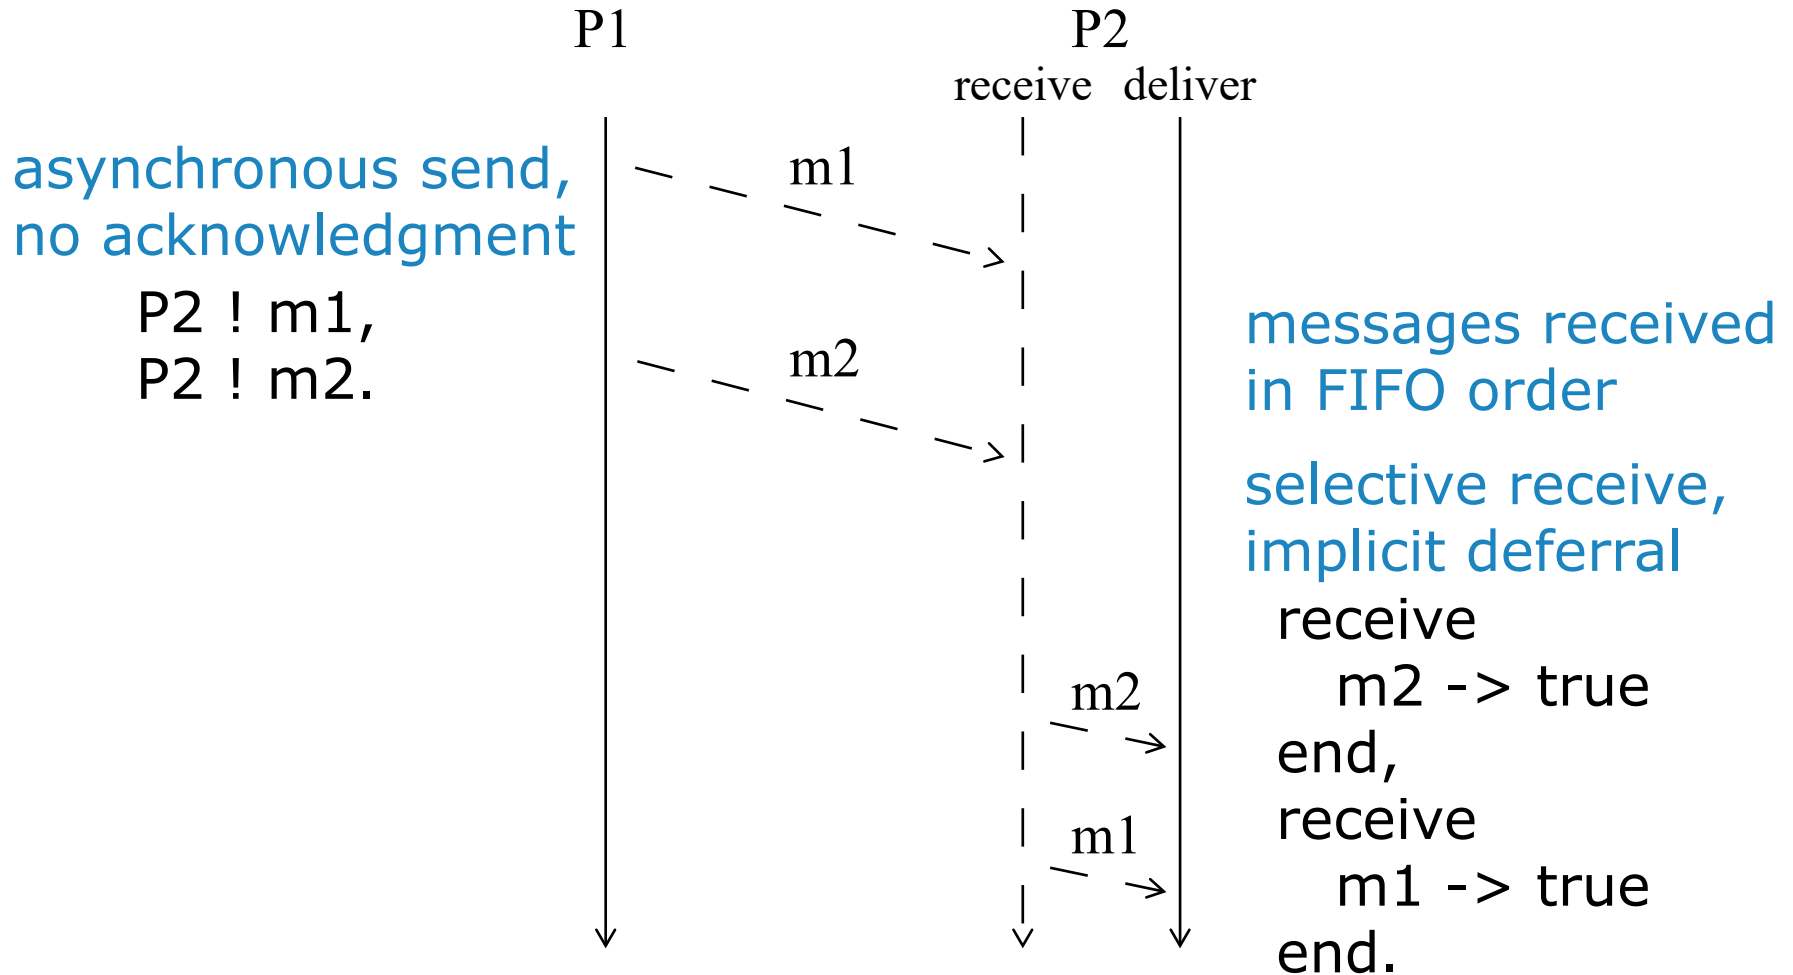
Pattern1 [when Guard1] -> Actions1 ;

... ;

PatternN [when GuardN] -> ActionsN

end.

# Message passing



# Message passing

---

## SENDER

P = spawn(wait, hello, []).

P ! "hello".

## RECEIVER

-module(wait).

-export([hello/0]).

hello() ->

receive

    X -> io:format("message received: ~s~n", [X])

end.

# Message passing

---

- Pids can be included in messages just like any data structure
  - Useful if we expect the receiver to reply

**SENDER**

Pid ! {ping, self()}.

**RECEIVER**

receive

{ping, FromId} -> FromId ! pong

end.

# Message passing

---

- receive can be augmented with a timeout
  - If no matching message has arrived within ExprT milliseconds, then ActionsT is evaluated instead

receive

```
Pattern1 [when Guard1] -> Actions1 ;  
... ;
```

```
PatternN [when GuardN] -> ActionsN
```

after

```
ExprT -> ActionsT
```

end.



# Registered processes

---

- register BIF allows giving names to processes  
register(Name, Pid)
  - Names can be used just as pids to send messages  
P = spawn(wait, hello, []).  
register(foo, P).  
foo ! "hello".
  - Name is automatically unregistered if the process terminates (use unregister/1 to do it manually)
  - Sending a message to a not registered name will cause a run-time error

# Links

---

- Two processes can be linked to each other
  - a) `link(Pid2)`
  - b) `Pid2 = spawn_link(Module, Func, Args)`
- Links are bidirectional and there can only be one link between two processes
- If one of the participants of a link terminates, it will send an exit signal to the other participant
  - Typically, if a process dies from an unexpected throw, error, or exit, the linked process also dies

# Monitors

---


- A process Pid1 can create a monitor for Pid2. The function returns a reference Ref
  - `Ref = erlang:monitor(process, Pid2)`
  - `{Pid2, Ref} = spawn_monitor(Module, Func, Args)`
- Monitors are unidirectional and repeated calls create several independent monitors
- If Pid2 terminates with exit reason Reason, a 'DOWN' message is sent to Pid1:
  - `{'DOWN', Ref, process, Pid2, Reason}`
- `erlang:demonitor(Ref)` BIF removes a monitor

---

# 0. Erlang


Distributed Programming

Facultat d'Informàtica de Barcelona (FIB)  
Universitat Politècnica de Catalunya (UPC)



# Distributed programming

---



- Each Erlang instance is a different node in the distributed system

Node 1: 'gold@host.upc.edu'

Node 2: 'silver@host.upc.edu'

# Distributed programming

---

- Make Erlang instances network aware
  - Set their name when starting the runtime via a command-line flag. Both long (-name) or short names (-sname) can be used (\*)

erl -name gold@host.upc.edu / erl -name gold

- node(): gold@host.upc.edu

erl -sname gold@host / erl -sname gold

- node(): gold@host

erl -name gold@127.0.0.1

(\*) A node with a long name cannot communicate with a node with a short name (and vice versa)

# Distributed process creation

---


- Create a process in a remote Erlang node
    - a) `P = spawn('gold@host.upc.edu', M, F, [A])`
    - b) `P = spawn('gold@host.upc.edu', fun (A) -> ... end)`
  - Connections are established automatically when another Erlang node is referenced
  - PIDs returned by `spawn` (or received in a message) can be used normally
    - Access transparent: local or remote are the same
    - Location transparent: process location is unknown
- P ! Message

# Registered processes

---

- Send a message to a process that is locally registered with a name on a remote node

{Name, Node} ! Message



{foo, 'gold@host.upc.edu'} ! "hello".

- Sending a message to a pid or name located at another node never fails



# Registered processes

---

- We will use locally registered names, but Erlang offers also globally registered names  
`global:register_name(foo, Pid).`
- Send a message to the process globally registered as `foo`  
`global:send(foo, "Hello").`
- Connections must be established explicitly  
`net_kernel:connect_node(Node).`

# Authentication

---

- If someone connects to a node, it gets connected to all the other nodes
- Use cookies as a mechanism to differentiate clusters of nodes  $\Rightarrow$  Nodes with different cookies are not able to communicate together
  - setcookie mycookie (command-line flag)
  - erlang:set\_cookie(node(), mycookie) (BIF)
  - Alternatively, you can have a file *.erlang.cookie* with the cookie in your home folder on all nodes

# More information

---

- Erlang official website: <https://www.erlang.org/>
- 'An Erlang Primer' by Johan Montelius
  - <https://people.kth.se/~johanmon/dse/crash.pdf>
- 'Learn You Some Erlang for Great Good!'
  - <https://learnyousomeerlang.com/>
- 'Concurrent Programming in Erlang, Part I'
  - <http://erlang.org/download/erlang-book-part1.pdf>
- Elixir language: <https://elixir-lang.org/>
  - Runs on the Erlang runtime